Autonomous Software Agents project

Bozzo Francesco, Izzo Federico

UNIVERSITÀ
DI TRENTO

2023-07-11

Contents

1 Agent

In the field of Computer Science, the term "agent" refers to an individual situated within an environment and capable of autonomously and flexibly taking actions to achieve its design objectives. Unlike traditional algorithms, agents do not require explicit definition of every edge case. Instead, a well-defined agent possesses a reasoning component that enables it to make decisions even in unknown situations. The flexibility of an agent can be assessed along two primary dimensions:

- Reactive: This dimension measures the delay required for the agent to respond to changes in the environment.
- Proactive: This dimension gauges the agent's ability to take proactive action to maximize future goals.

Communication models between agents and their environments vary depending on the specific characteristics of the environment and the agent itself. Generally, an agent perceives observations from the environment through sensors and carries out actions on the environment using actuators.

Autonomy is a fundamental characteristic of an agent. The internal decision-making process of an agent, often referred to as its "brain," should be capable of handling decisions with or without collected information. Furthermore, it should be able to adapt and evolve in response to potential changes in requirements.

Agents can be designed to solve tasks or goals. Task-oriented agents focus on accomplishing smaller objectives that contribute to the achievement of a larger final goal. On the other hand, goal-oriented agents receive a specific goal and autonomously determine a list of tasks necessary to fulfill the assigned goal.

1.1 Multi-agent system

A multi-agent system, as implied by its name, refers to a collection of agents situated within the same environment. Interactions within such a system can be broadly categorized into two types: cooperative and competitive. Throughout this project, we will delve into both of these interaction modes at various stages of development.

In a competitive system, multiple agents act in opposition to one another, where the overarching goal can be divided into sub-goals focused on maximizing personal rewards while minimizing opponents' gains. Within this scenario, the objective function may be shared among enemy agents, and it is also plausible to have multiple functions where each agent interferes with enemies solely to achieve its own goals.

On the other hand, a cooperative system consists of numerous agents working together to maximize a shared reward. In such systems, each agent must possess the capability to cooperate, coordinate, and engage in negotiation to the greatest extent possible. Cooperative systems can be further classified into two main categories:

- Simple (reciprocal) cooperation: This form of cooperation occurs when the benefits derived from collaboration outweigh the costs associated with the actions taken. It is considered the simplest type of cooperation as it leads to increased fitness for both the helper and the helped parties.
- Altruistic cooperation: In this case, the cost incurred by the individuals or species offering assistance surpasses the advantages gained. This approach is often regarded as more challenging since it cannot be readily explained by a purely "genetic-centric" perspective.

An essential aspect to consider when implementing a cooperative system is the communication mechanism. It should prioritize speed, reliability, and minimize delays as much as possible.

1.2 Architecture

There are several architectural options available for constructing an agent with the ability to operate within a specific environment. For our purposes, we have chosen to adopt the architecture outlined in the following pseudocode.

Algorithm 1 Agent control loop

 1: **procedure** AGENTCONTROLLOOP
 2:     $B \leftarrow B_0$ ▷ Belief set initialization
 3:     $I \leftarrow I_0$ ▷ Intention set initialization
 4:     **while** true **do**
 5:         perceive ρ
 6:         $B \leftarrow \mathsf{update}(B, \rho)$ ▷ Belief set update
 7:         $D \leftarrow \mathsf{options}(B)$ ▷ Desires computation
 8:         $I \leftarrow \mathsf{filters}(B, D, I)$ ▷ Intention update
 9:         $\pi \leftarrow \mathsf{plan}(B, I)$ ▷ Plan computation
10:         **while** not (empty(π) or succeeded(I, B) or impossible(I, B)) **do**
11:             $\alpha \leftarrow hd(\pi)$ ▷ Get next action
12:             execute(α) ▷ Execute action
13:             $\pi \leftarrow \mathsf{tail}(\pi)$ ▷ Remove executed action
14:             perceive ρ
15:             $B \leftarrow \mathsf{update}(B, \rho)$ ▷ Belief set update
16:             **if** reconsider(I, B) **then**
17:                 $D \leftarrow \mathsf{options}(B)$ ▷ Desires computation
18:                 $I \leftarrow \mathsf{filters}(B, D, I)$ ▷ Intention update
19:             **end if**
20:             **if** not sound(π, I, B) **then**
21:                 $\pi \leftarrow \mathsf{plan}(B, I)$ ▷ Plan computation
22:             **end if**
23:         **end while**
24:     **end while**
25: **end procedure**

After each communication with the environment, the belief set, denoted as B, undergoes an update. Subsequently, the desires set, referred to as D, is computed based on the updated belief set and then filtered to generate the intentions set, denoted as I. The intentions set, in conjunction with the belief set, is utilized to formulate a comprehensive plan, denoted as π, which aims to fulfill all the intentions. This plan is then executed incrementally, action by action, represented as α. Whenever the agent receives new information from the environment, the belief set, intentions set, and desires set are updated accordingly. Following this update, the agent can decide whether to generate a new plan or adhere to the existing one.

2 Deliveroo

The Deliveroo environment serves as the designated playground for developing our BDI agents. It encompasses a delivery-based game, where agents navigate on a two-dimensional grid. The objective of the game is to collect parcels dispersed throughout the map and swiftly deliver them to designated delivery zones. Each parcel possesses an assigned reward value, which may decay over time, and is granted to the agent responsible for its successful delivery.


Figure 1: An example of a Deliveroo map.

The game can be played with either a single agent or multiple agents. In the latter case, agents are solid entities capable of obstructing each other's movement on the road.

Agents typically have limited perception of the world, restricting their awareness to other agents and parcels within a specified radius. Additionally, agents have the ability to carry and deliver multiple parcels simultaneously.

Deliveroo has been designed to offer diverse game scenarios through the manipulation of various parameters, including:

- Parcel: Generation interval, reward decay, and reward value distribution.
- Player: Number of steps and movement speed.
- External agent: Quantity and movement speed.
- Sensing: Radius of parcel and agent perception.
- Map: Definition of the grid board and classification of tiles as walkable, non-walkable, or delivery zones.
- Game clock.

Specifically, this examination project comprises two distinct deliverables, each with a different solution strategy:

- Single-agent: The agent operates independently, aiming to maximize its individual score.
- Multi-agent: The agent engages in communication with other agents to exchange information and collectively optimize the overall score of the group.

3 Single agent implementation

Considering the intricate nature of the environment, a multitude of tailored solutions have been implemented to address various complexities, difficulties, and challenges encountered.

3.1 Manhattan distance

The Manhattan distance is a mathematical function that offers a rapid estimation of the distance between two points. It is calculated by summing the absolute differences between the respective x-coordinates and y-coordinates of the two points:

$$d_{ab} = |b_x - a_x| + |b_y - a_y|$$

While this function may not be fully suitable for the Deliveroo game due to the presence of non-walkable tiles on the maps, it serves as a reasonable and computationally efficient heuristic for approximating distances in certain cases.

3.2 Problem parameters estimation

As detailed in Section 2, various parameters can be adjusted during the initialization of the game. While certain parameters are directly communicated to the agent, others remain obscure and can only be estimated. To enhance the accuracy of estimating player rewards, we have implemented a learning mechanism aimed at approximating the player's speed and the decay of parcel rewards over time. This mechanism enables us to build a more precise estimation for optimizing player rewards.

3.2.1 Player speed

During the agent's movement within the map, it maintains a record of its position over time. This historical positional data, along with corresponding timestamps, allows the agent to estimate its own speed using the following algorithm:

Algorithm 2 Player speed estimation

1: **procedure** UPDATEMAINPLAYERSPEEDESTIMATION($\mathcal{D}, s, \phi$)
2:      deltas $\leftarrow []$ ▷ Initialize an empty array of deltas
3:      **for each** timestamp $t_i \in \mathcal{D}$ **do**
4:          $deltas.append(t_i - t_{i-1})$ ▷ Delta of two consecutive timestamps
5:      **end for**
6:      $c \leftarrow s * (1 - \phi)$ ▷ Current speed contribution
7:      $n \leftarrow \text{avg}(deltas) * \phi$ ▷ New speed contribution
8:      $s \leftarrow c + n$ ▷ New estimation
9:      **return** c
10: **end procedure**

Here, ϕ represents the learning rate hyper-parameter, which can be utilized to control the influence of the most recent measured instantaneous speed in relation to its historical value.

3.2.2 Parcel decay

Apart from estimating the agent's speed, our agents are also capable of estimating the decay of parcel rewards over time. Similar to the player speed estimation, the parcel decay is calculated based on differences in timestamps. Each timestamp is associated with a sensed reward update of a visible parcel, allowing us to estimate the decay of the parcel rewards as time progresses.

Algorithm 3 Get parcel decay estimation

1: **procedure** GETPARCELDECAYESTIMATION($\mathcal{D}$)
2:      deltas $\leftarrow$ []   ▷ Initialize an empty array of deltas
3:      **for each** timestamp $t_i \in \mathcal{D}$ **do**
4:          $deltas.append(t_i - t_{i-1})$ ▷ Delta of two consecutive timestamps
5:      **end for**
6:      **return** deltas
7: **end procedure**

Algorithm 4 Parcels decay estimation

1: **procedure** UPDATEPARCELSDECAYESTIMATION($\mathcal{P}, d, \phi_2$)
2:      deltas $\leftarrow$ []   ▷ Initialize an empty array of deltas
3:      **for each** parcel $p_i \in \mathcal{P}$ **do**
4:          $deltas.concat(getParcelDecayEstimation(p.timestamps))$ ▷
5:      **end for**
6:      $c \leftarrow d * (1 - \phi_2)$ ▷ Current decay contribution
7:      $n \leftarrow \text{avg}(deltas) * \phi_2$ ▷ New decay contribution
8:      $d \leftarrow c + n$ ▷ New estimation
9:      **return** d
10: **end procedure**

Here, ϕ_2 represents the learning rate, which controls the contribution of past estimations relative to the current estimated parcel decay. It allows us to regulate the influence of previous estimations when updating and refining the estimation of the parcel decay over time.

3.3 Probabilisitic model

Within the environment, multiple competitive agents coexist, and their effectiveness in picking up parcels significantly impacts the value of each individual parcel. To address this, we have developed a penalty value based on a probabilistic model that takes into account the potential plans of other competing agents.

The underlying concept behind this probabilistic model can be summarized as follows: "If there is a parcel available and I am the closest agent to it, I have a higher probability of reaching and acquiring it faster than any other agents. Consequently, this parcel should be given more weight and consideration, even if its assigned value is lower than that of other parcels located further away." This assertion can be formulated more formally as follows:

$$\text{penalty probability} = \frac{\sum_{a \in \mathcal{A}} \frac{d_{max} - d_{pa}}{d_{max}}}{|A|}$$

Here, we denote $\mathcal{A}$ as the set of opponent agents, d_{max} as the maximum distance between the parcel and the collective group comprising opponent agents, the main player, and cooperative agents. Additionally, d_{pa} represents the distance between the parcel and an opponent agent.

3.4 Potential parcel score

The process of parcel selection plays a crucial role in defining an effective agent. To accurately estimate the potential reward gain of a parcel, our agents consider various elements and metrics. Formally, the final reward for a parcel is computed as follows:

$$r_f = r - \left(d_{ap} * \frac{s_a}{decay} \right) - \left(d_{min} * \frac{s_a}{decay} \right) - r * \text{penalty probability}$$

Here, d_{ap} represents the distance between the agent and the parcel, s_a denotes the estimated speed of the agent, d_{min} represents the minimum distance between the parcel and the nearest delivery zone, and penalty probability corresponds to the probability calculated using the probabilistic model discussed in Section 3.3.

The resulting formula takes into consideration multiple factors, including:

- The reward that the agent expends to approach the parcel and deliver it to the nearest delivery zone, which represents the minimum cost associated with delivering that particular parcel.
- An approximate estimation of other agents' intentions based on their distances from the parcels, incorporating a probabilistic model.

3.5 Distances cache

In order to optimize computational efficiency and enhance the accuracy of reward estimation, a cache is maintained to store distances between tiles throughout the map. Whenever a plan is generated, the distance between the starting point and any other tile along the path is stored in the cache. However, it should be noted that this approach does not guarantee the shortest route between two tiles. As a result, cache entries are updated whenever a smaller distance value is discovered. This caching mechanism acts as a form of learning, gradually improving over time.

Throughout the codebase, the cache is utilized in numerous instances. In the event of a cache miss, the agent resorts to using the Manhattan distance as a fallback measure. By employing this caching strategy, the goal is to strike a balance between computation efficiency and accurate reward estimation.

3.6 Replan

As Deliveroo is a dynamic game that involves simultaneous actions from multiple agents, we have implemented a mechanism to replan the actions of the current agent if it fails to execute a move within a specific time frame. This functionality allows agents to prevent getting stuck in narrow or crowded areas of the map. By triggering a replanning process when necessary, agents can adapt their actions and navigate through challenging situations more effectively.

4 Multi agent implementation

The communication protocol utilized in our implementation is built upon a library that provides various endpoints to facilitate the handling of different types of messages. These endpoints include "say" for regular communication, "shout" for broadcasting messages to all agents, "ask" for querying other agents, and "broadcast" for widespread information dissemination. By leveraging these endpoints, agents are able to engage in effective communication and exchange relevant information during the game.

4.1 Information sharing agents

During the game, teams have the ability to share sensed data from the environment among their members. For simplicity, our agents utilize broadcast messages as the means of sharing information, making them visible to all connected agents in the game.

This allows each agent to construct its own plan based on the information sensed by all other agents across the map. Additionally, agents communicate their intended parcel pickups and the corresponding plans via broadcast messages. This mechanism helps reduce collisions among agents on the map and prevents unnecessary detours to pick up parcels that have already been claimed.


Figure 2: Information sharing.

By sharing information, agents are able to compute more refined plans that take into account hidden map locations. This multi-agent implementation follows a distributed model that can scale effectively with a large number of agents. Moreover, since there is no central computing unit, agents can be added or removed at any time.

It is important to note that this approach relies on broadcast communication, which means that any malicious agent could potentially understand the communication protocol and inject false information. To address

this issue, we have considered implementing an encryption mechanism based on an initial exchange of keys. However, considering the nature of the course, we opted for a simpler implementation, focusing on other important implementation details.

4.2 Leader-members agents

The leader negotiation process plays a pivotal role in the multi-agent architecture, whereby one agent is designated as the leader responsible for computing plans for all other agents.

4.2.1 Leader negotiation

The process of electing a leader is a well-known problem in computer science, but for the purpose of our project, we opted for a simple solution due to our focus on other aspects. The negotiation for the leader role is facilitated through two types of messages: "askforaleader" and "leader".

Upon connecting to the game and receiving their initial position, each agent broadcasts an "askforaleader" message, inquiring if a leader has already been elected. This message serves as a request for information regarding the existence of a leader. Conversely, the "leader" message is used by the agent who has been elected as the leader to communicate its identity.

Following the transmission of the "askforaleader" message, a timeout period of 2.5 seconds is set. If no response from an existing leader is received within this timeframe, it indicates that no leader has been elected yet. In such a case, the agent who sent the "askforaleader" message assumes the role of the leader and broadcasts a message to inform other agents of its newly elected leader status.


Figure 3: Leader negotation.

It is important to note that the simple implementation described above may introduce rare scenarios where multiple agents connect at exactly the same time, potentially leading to a race condition and the presence of multiple active leaders.

This enhanced model implementation can be considered an extension of the approach discussed in Section 4.1. In addition to sharing information about non-visible areas, this model allows for more comprehensive decision-making by leveraging knowledge derived from the computation of all plans. However, it comes with certain drawbacks. Firstly, it introduces a single point of failure, as all plans are generated by a single node. Secondly, the scalability of the system is limited when dealing with a large number of nodes. On the positive side, the system offers enhanced security, as plan communication occurs through point-to-point communication channels that cannot be accessed by malicious agents.

4.2.2 Plan communication

The plan communication system operates in a straightforward manner. When an agent does not have a plan, it initiates a request for a new plan by sending an askforplan message. This request is implemented using a point-to-point ask primitive, which ensures that the request is directed specifically to the leader. Upon receiving the request, the leader begins the process of computing the plan. Once the computation is complete, the leader sends the list of actions comprising the plan back to the original agent using another point-to-point communication. This ensures that the plan is securely and efficiently transmitted between the leader and the requesting agent.


Figure 4: Plan communication.

4.2.3 Traffic penalty

In this communication model, the leader serves as the central compute node responsible for generating plans for all agents. This central position grants the leader extensive knowledge about the future movements of other agents. To enhance the computation of potential parcel scores described in Section 3.4, we have introduced an additional penalty that accounts for traffic considerations and aims to create plans that evenly distribute agents across the entire map.

To facilitate this, a traffic map is maintained on the leader. This map is a copy of the original map, and it is updated every time a plan is generated. For each tile included in a plan, the corresponding position on the traffic map is incremented by 1. When an agent requests a new plan, the previously computed plan for that agent is used to decrement the corresponding positions on the traffic map. This approach allows for the consideration of traffic patterns and encourages agents to choose paths that minimize congestion and evenly distribute their movements throughout the map.


Figure 5: Empty traffic map to traffic map with one plan.

In Figure 5 it is presented an empty traffic map (on the left) with the respective delivery zones, after the computation of a plan the traffic map is updated (on the right). This process is applied for every generated plan and the final result is presented in Figure 6, it is clear that some tiles are more trafficate than others.


Figure 6: Traffic map with multiple plans.

With the traffic map it is possible to take a parcel and analyze its neighbours in order to understand if it is a trafficated area and consequentially if it is a good idea to take it.

The logic behind the traffic penalty is summarized in the following pseudocode:

Algorithm 5 Traffic penalty

1: **procedure** TRAFFICPENALTY(M, p)
2:     $m \leftarrow \max(M)$ ▷ Obtain the maximum traffic in the current traffic map
3:     $t \leftarrow 0$ ▷ Initialize neighbourhood traffic
4:     $ns \leftarrow \text{getNeighbours}(p)$ ▷ Get parcel neighbour tiles
5:     **for each** neighbour $n_i \in ns$ **do**
6:         $traffic \leftarrow traffic + M[p.x][p.y]$ ▷ Update neighbourhood traffic
7:     **end for**
8:     $t \leftarrow t/\text{len}(ns)$ ▷ Average neighbourhood traffic
9:     $p \leftarrow \min(t/m)$ ▷ Obtain a probability from average traffic
10:     **return** $2 * p.reward * p$
11: **end procedure**

4.3 Implementation comparisons

Distributed based	Leader based
Malicious information injection	Single point of failure
Scalable on the number of agents	Single compute node
Resilient	Ability to compute traffic map

5 Benchmarking

6 Conclusion

7 References